

# Z-Journal: Scalable Per-Core Journaling

Jongseok Kim<sup>1</sup>, Cassiano Campes<sup>1</sup>,  
Joo-Young Hwang<sup>2</sup>, Jinkyu Jeong<sup>1</sup> and Euseong Seo<sup>1</sup>



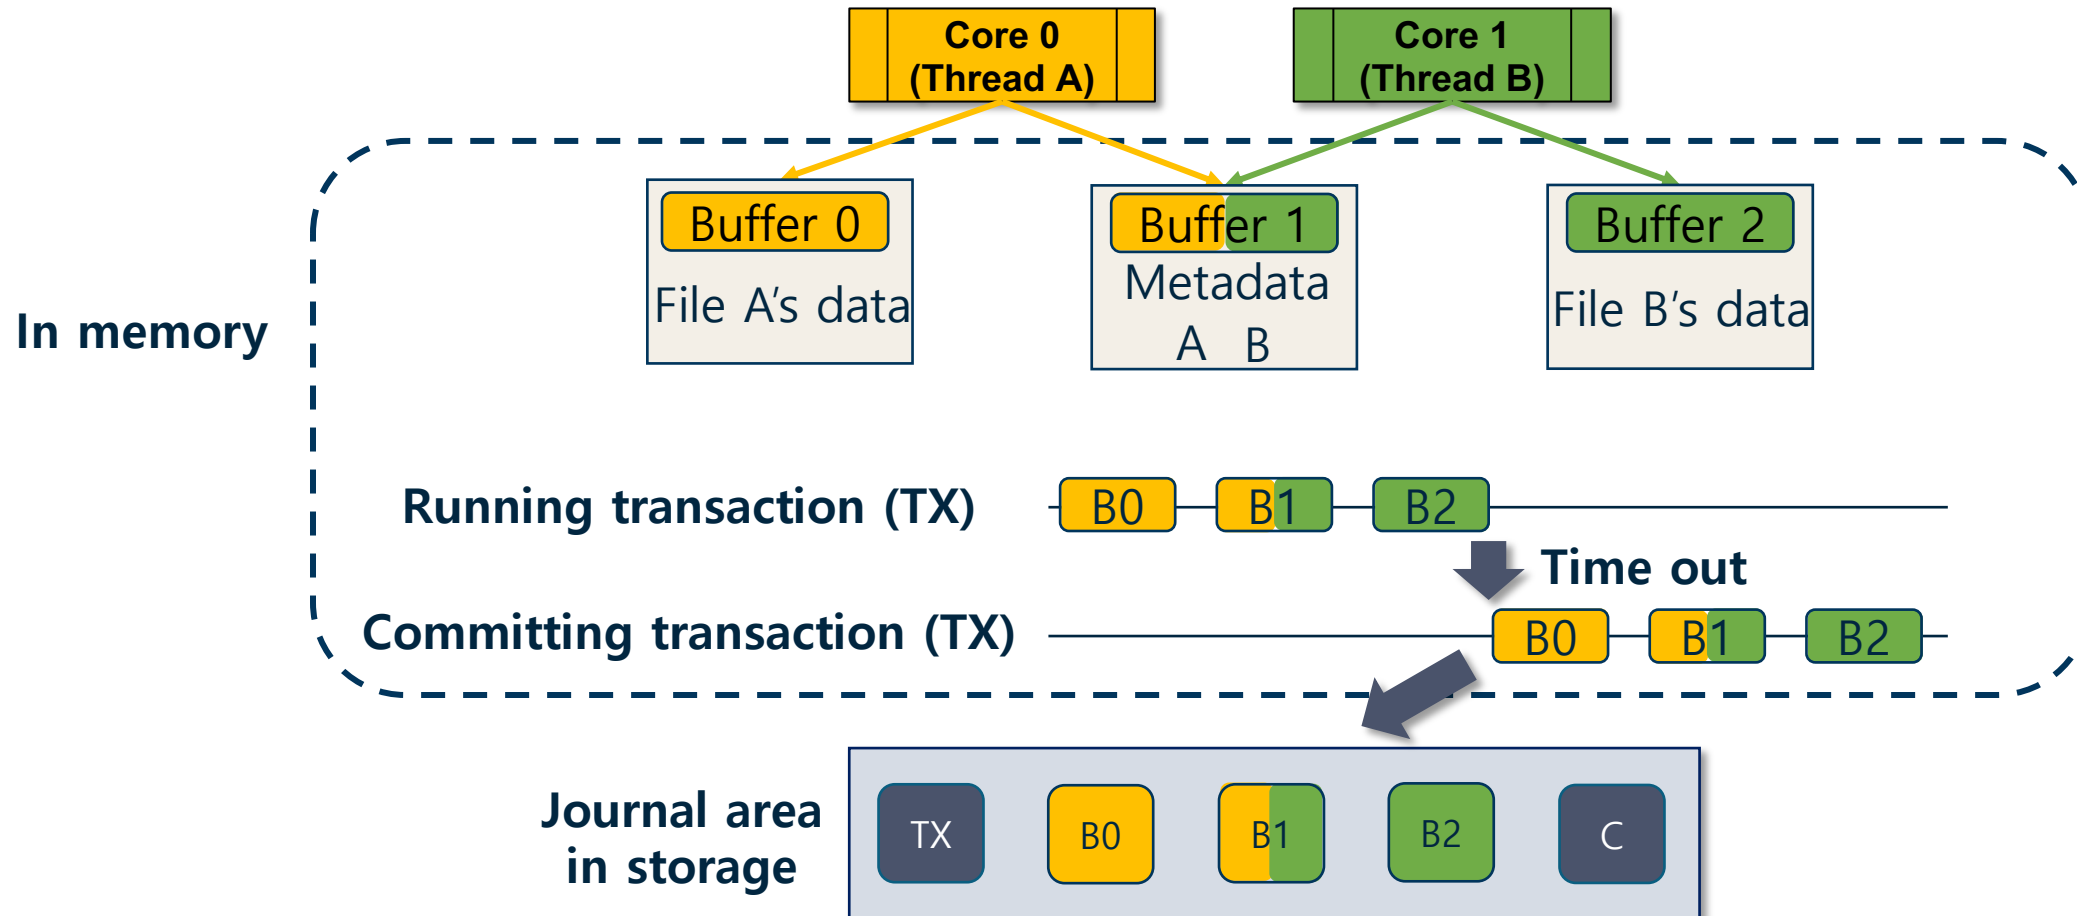
Sungkyunkwan University<sup>1</sup>



Samsung Electronics Co., Ltd.<sup>2</sup>

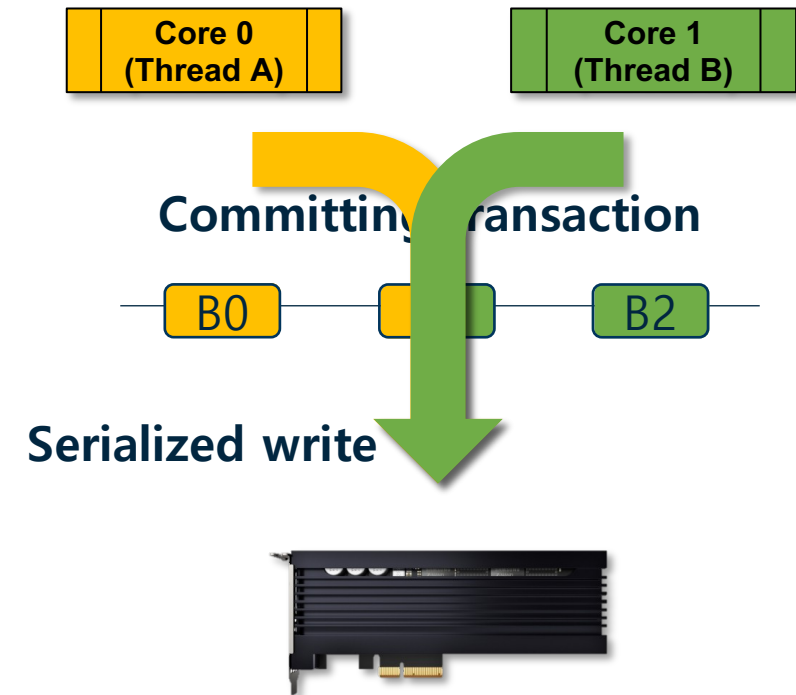
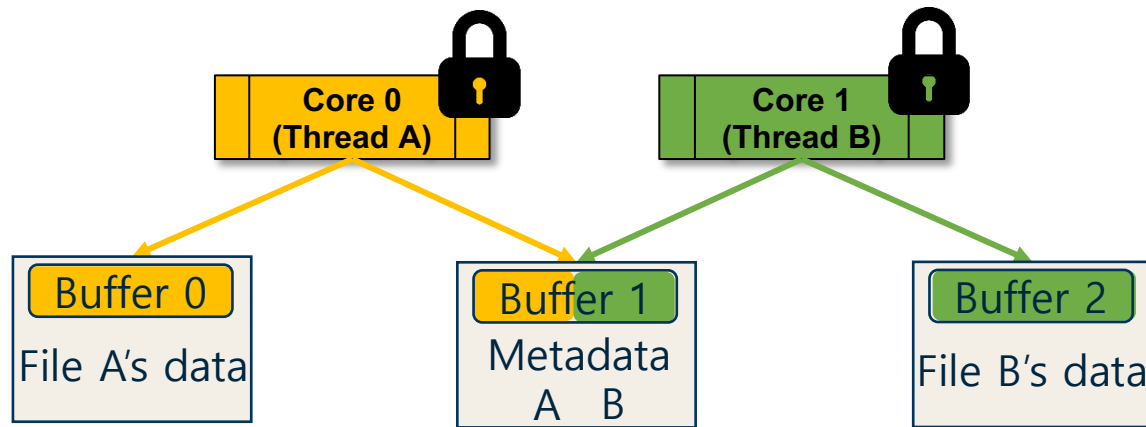
# Conventional Journaling Scheme

- JBD2: Generic Journal Layer for Ext4 and Other FS



# Journal – A Scalability Bottleneck

- The **lock acquisition** for the running transaction
- Not to fully **utilize the internal-parallelism** provided by the modern NVMe SSDs



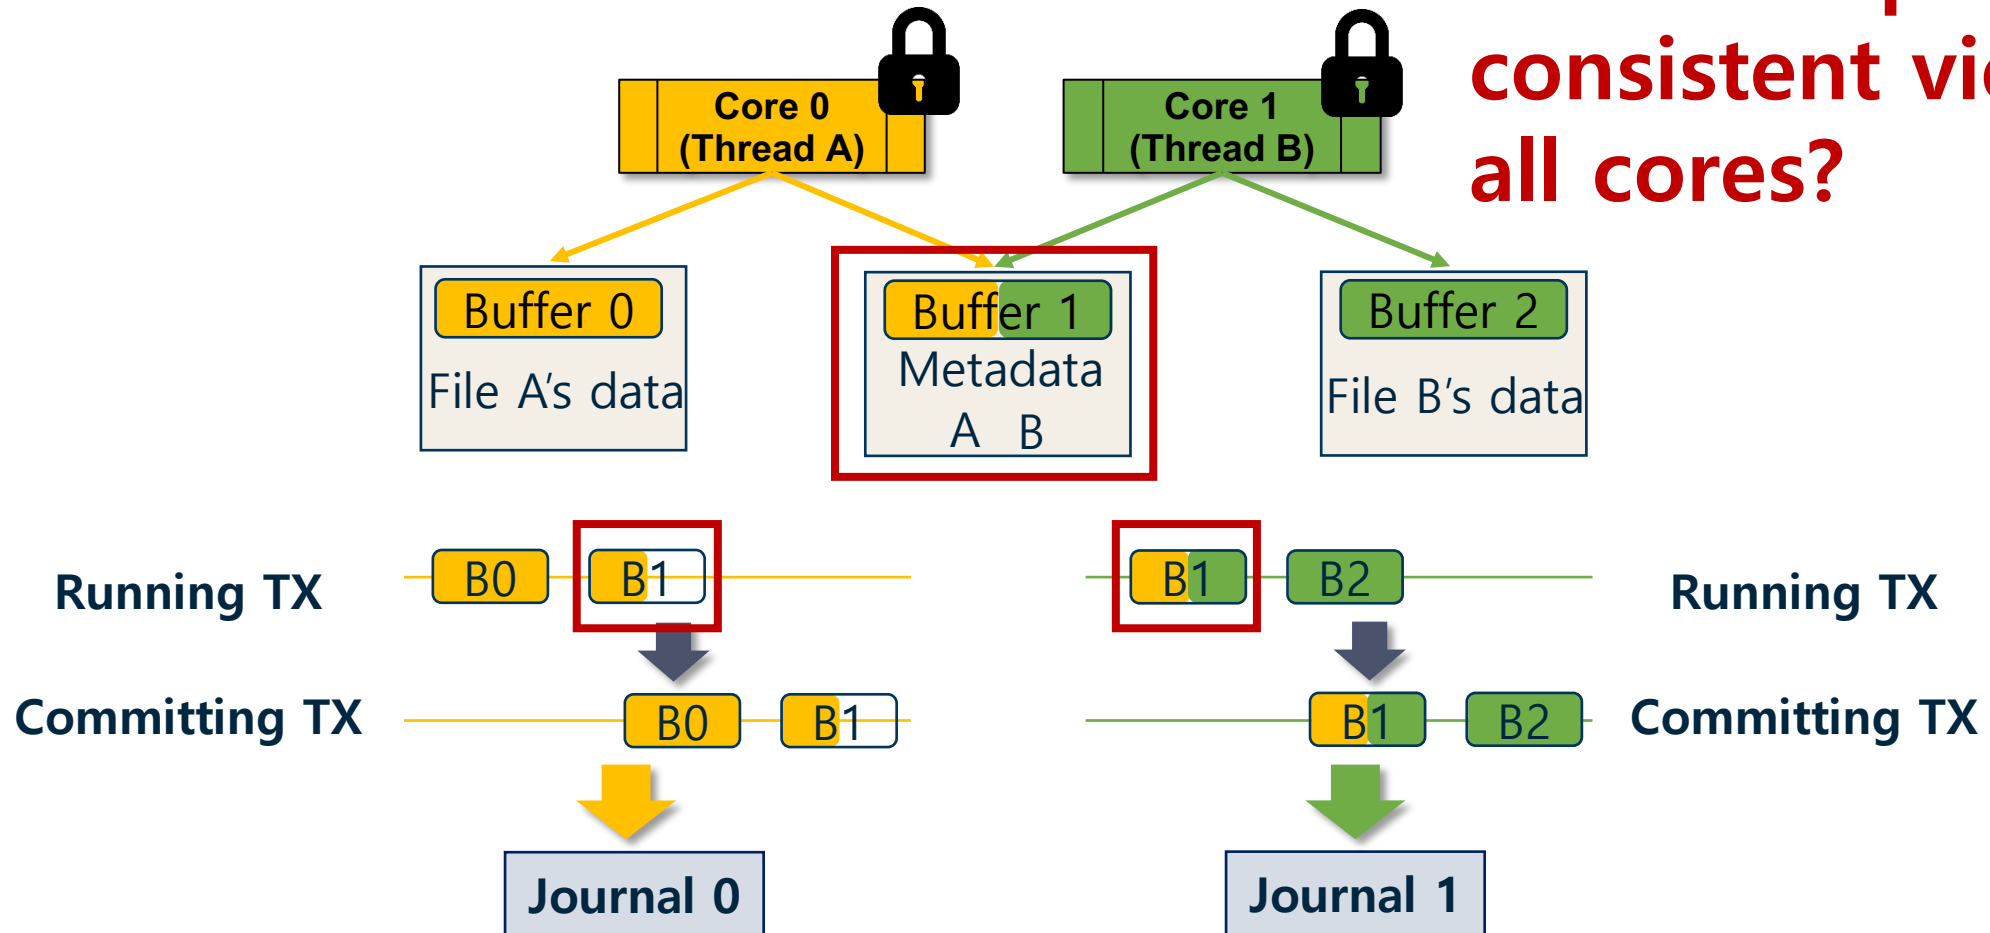
# Related work

- **Parallelize a part of journal layer**
  - iJournaling [ATC '17]
  - High-performance transaction processing [FAST '18]
  - ➔ **Inherent serialization from committing to a centralized journal remains**
- **Fully-redesigned file systems for scalability**
  - SpanFS [ATC '15]
  - ScaleFS [SOSP '17]
  - ➔ **Journal layer is tightly coupled with file system layer**  
**They cannot be directly applied to the existing file systems**

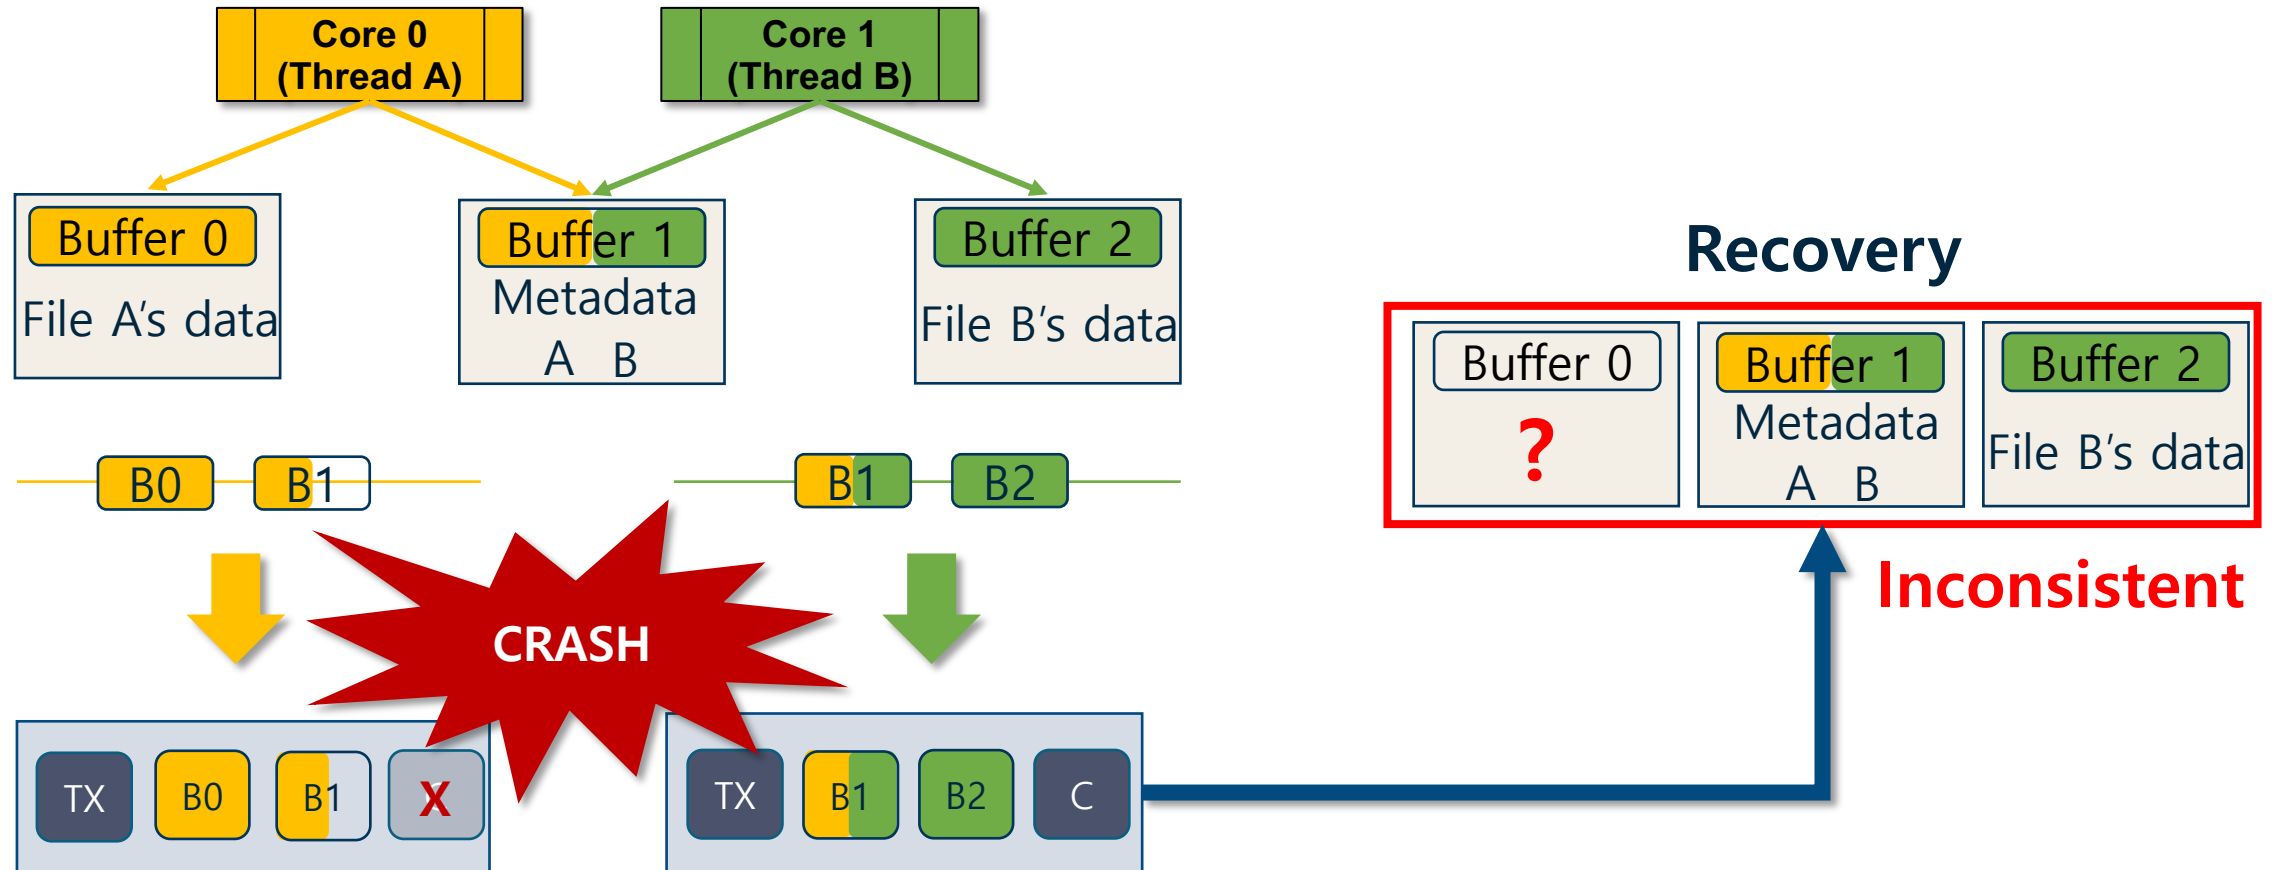
# Intuitive Solution

- Per-core Journal

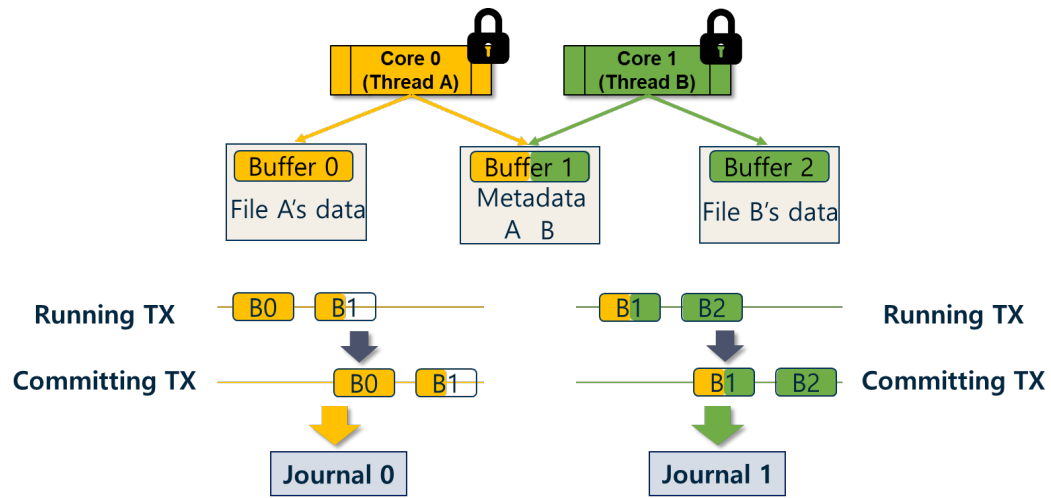
Does this provide  
consistent view to  
all cores?



# Journal Coherence Problem



# Our Approach: Z-Journal



Per-core Journaling



- Journal coherence commit
  - Order preserving transaction chaining
- Proactive frozen copy
- Core metadata allocation
- Please refer to our paper
- Journal coherence checkpointing
- Recovery

Journal coherence mechanism

# Order-preserving Transaction Chaining

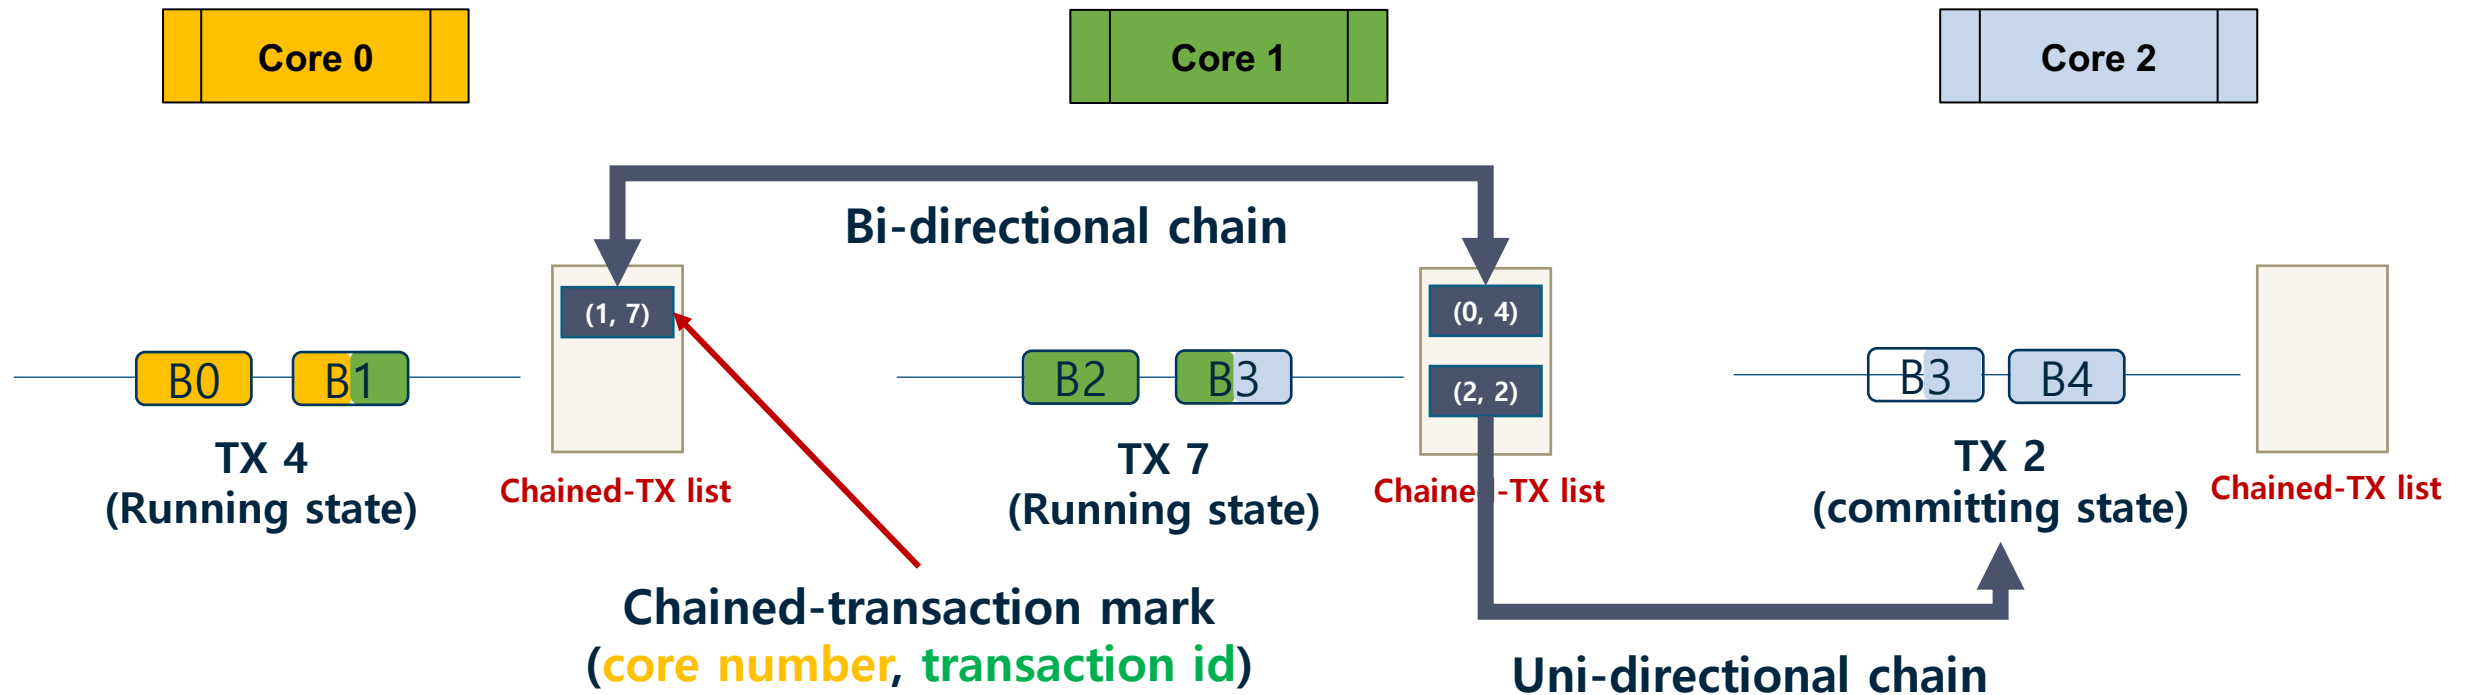
- We can checkpoint a modification of a buffer only after its preceding modifications become persistent
- We need to record write orders to a shared block
- The transaction chain graph
  - Imposing order-constrains over transactions
  - Putting off the enforcement of write-order constraints to the checkpoint and recovery time



# Construction of Transaction Chain Graphs

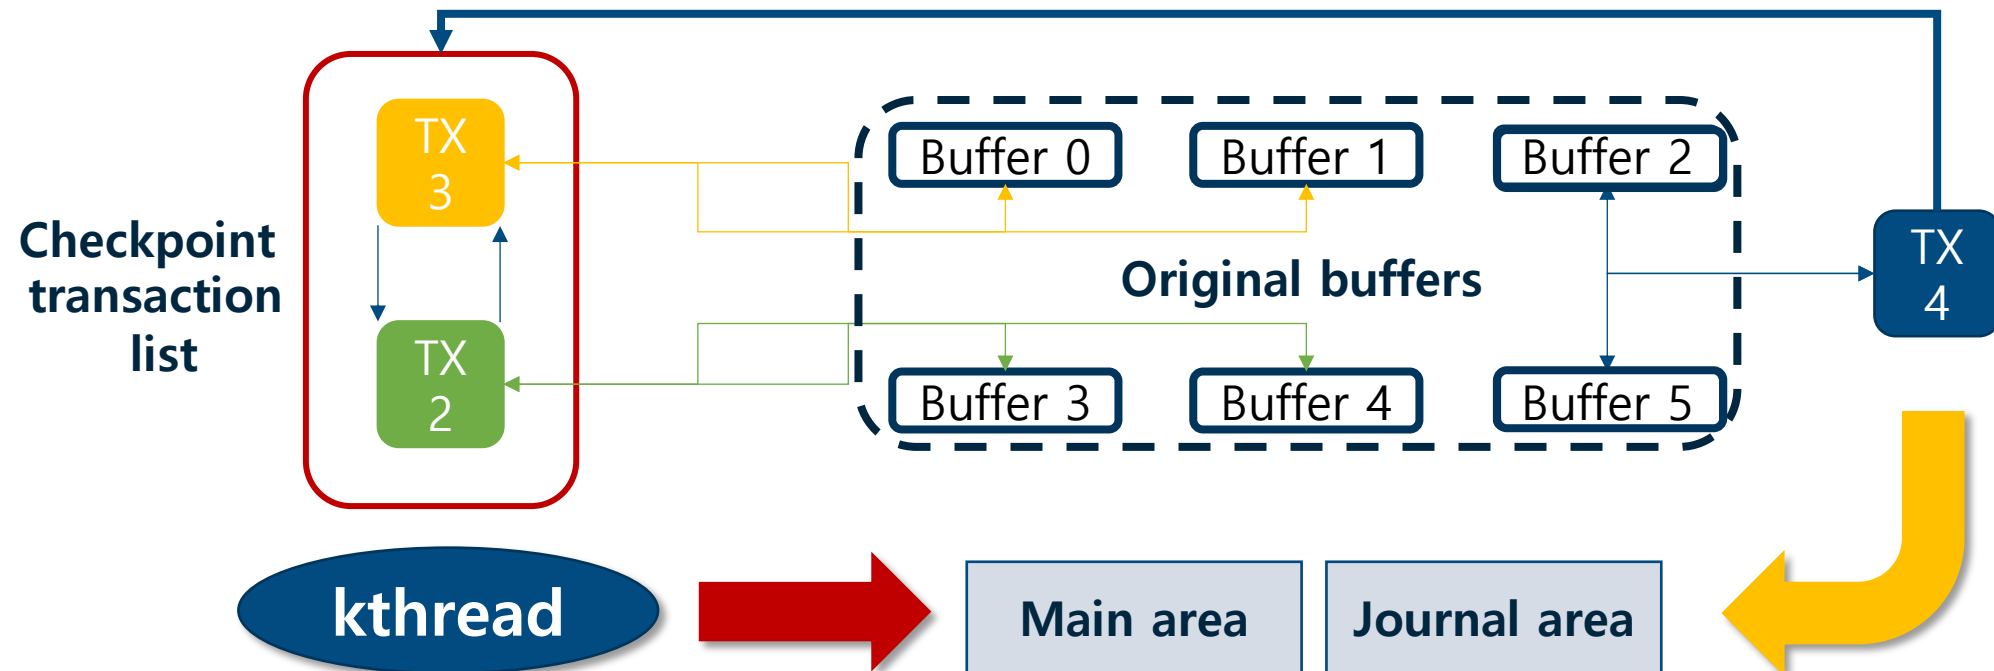
- Chained-transaction list
- Chained-transaction mark

Allowing each core to **commit** transactions to its journal **independently** to other journals



# Conventional Checkpoint Process

- **When the commit is finished**
  - The transaction is inserted at the checkpoint transaction list
  - The buffers are **marked as dirty**



# Journal Coherence Checkpoint


- **When the commit is finished**
  - The transaction will be **skipped over setting dirty flags** of transaction's buffers
  - As stated earlier, not all committed transactions become objects of checkpointing in Z-Journal
- **A transaction is *valid***
  - A transaction without a preceding transaction is *valid*
  - When all of its preceding transactions are *valid*
- **Z-Journal checkpoints only the valid transactions**
  - If a transaction turns out to be valid, its buffers will be marked as dirty

# Transaction Chain Graph Traversing

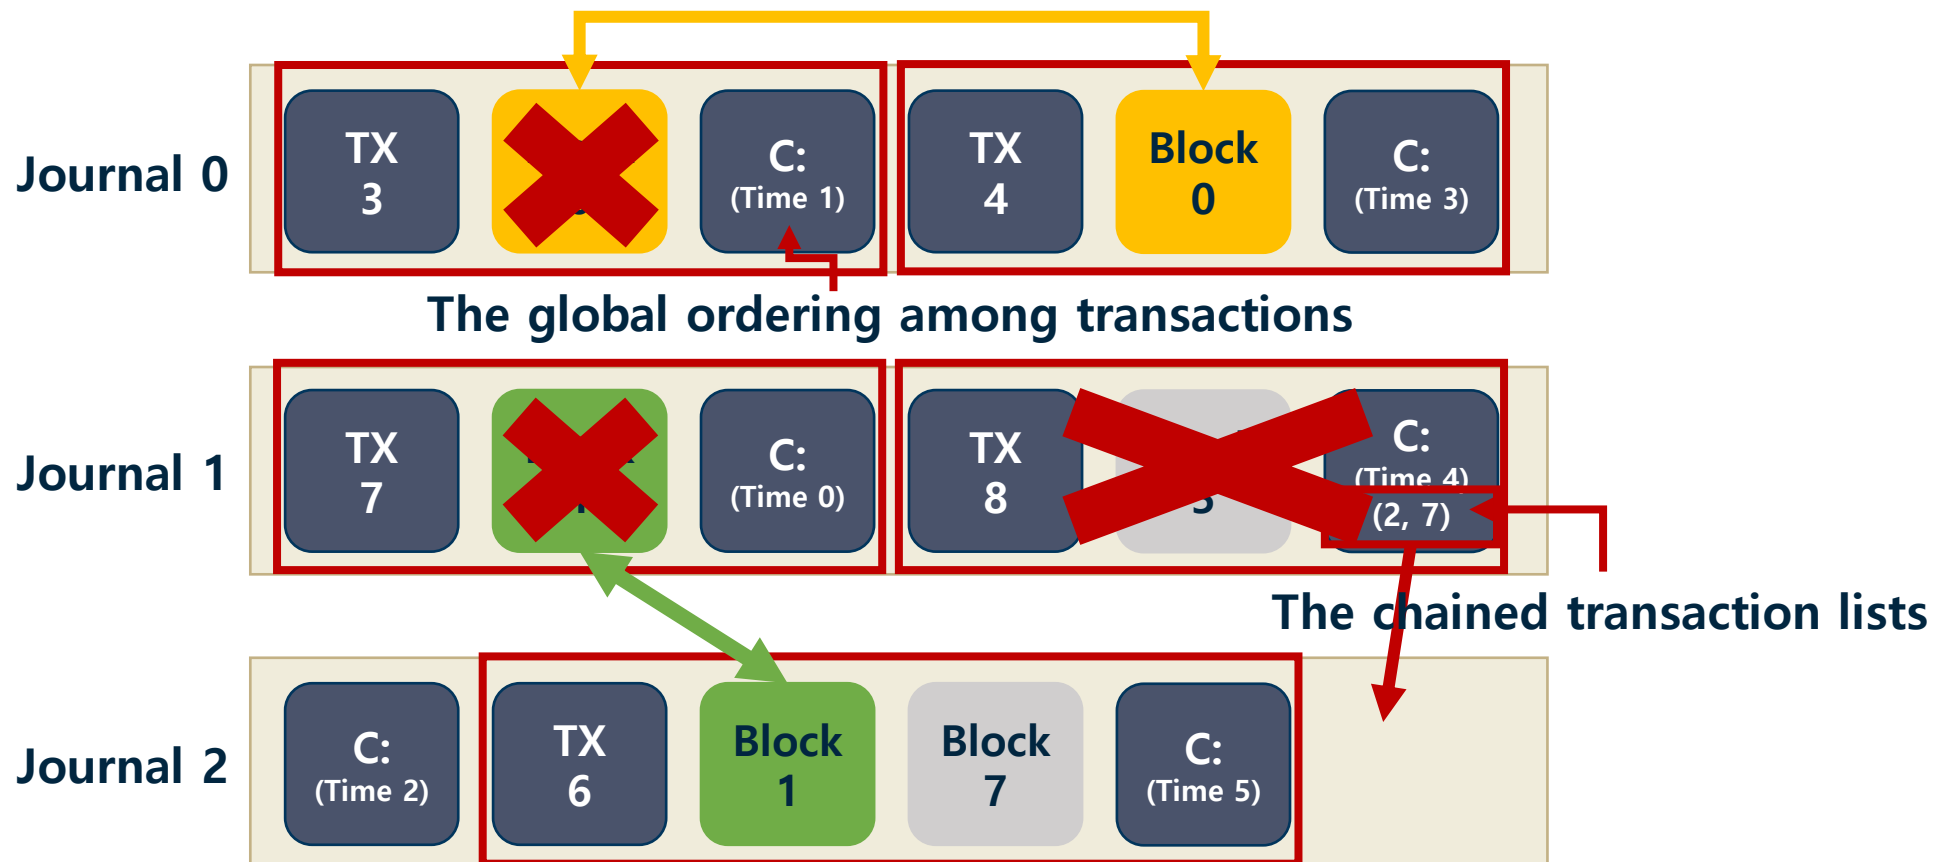
- To check the validity of a committed transaction, traversing the transaction chain graph is required

- **Examples**

- Check TX 2
- Check TX 3
- Check TX 4



# Recovery



- Traversing **the graphs** to find valid transactions
- Comparing **the timestamps** of both transactions to determine the latest buffer image to restore

# Evaluation

- **Environment**

	Specification
Processor	Intel Xeon Gold 6138 x 4 sockets
Memory	DDR4 2666 MHz 32GB x 16
Storage	Samsung SZ985 NVMe SSD 800GB
OS	Linux kernel 4.14.78 (Journal=data mode journaling)

- **Implementation**


- Z-Journal is modified based on [JBD2](#)
- The [ext4 file system](#) has minimal modifications to recognize multiple journals

- **Workloads**


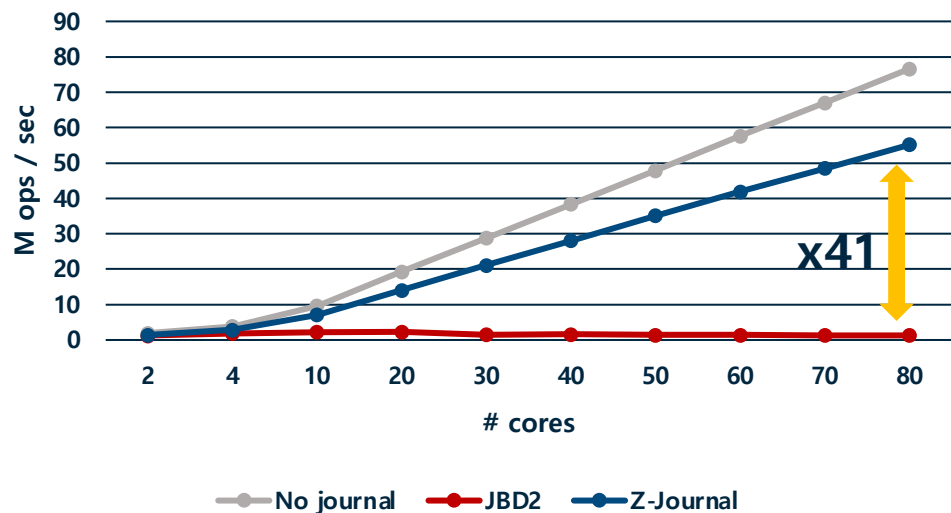
- FxMark<sup>1</sup>, Filebench, SysBench

<sup>1</sup> Changwoo Min, Sanidhya Kashyap, Steffen Maass, and Taesoo Kim. *Understanding manycore scalability of file systems*. In 2016 USENIX Annual Technical Conference (ATC), 2016.


# Under Various Sharing Conditions



Low sharing




Medium sharing



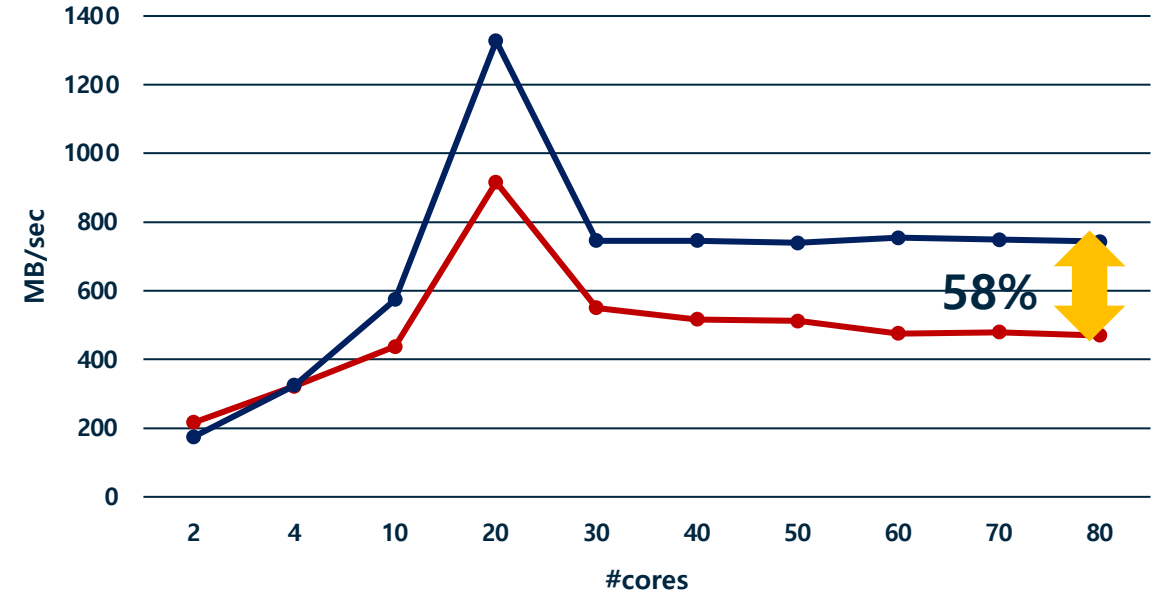
Overwrite

# Overall File System Performance



—●— JBD2 —●— Z-Journal

**Fileserver**



—●— JBD2 —●— Z-Journal

**Varmail**



# Z-Journal Conclusion

- **Per-core journal**
  - The thread running on a core can write to **the journal dedicated to the core independently** to the other threads
- **Journal coherence mechanism**
  - This enables scalable per-core journaling while **keeping crash consistency**
- **Performance evaluation**
  - Our evaluation showed that, through its per-core journaling design, Z-Journal is faster and more scalable than the current JBD2

**Thank you**

**ks7sj@skku.edu**